



Decomposing Autonomous Pathogen Creation: A Graph-Based Safety Evaluation Framework

Zuzanna Matuszewska

1. Threat Model and Theory of Change

Threat Model


- Autonomous AI agents + synthetic biology advances are reshaping the biosecurity landscape.
- Capable systems compress expertise, coordinate complex steps, and scale effort.
- New risk: misaligned AI could independently pursue bioweapon development.

BARRIER LOWERING Novice + LLM	CEILING RAISING Expert + AI Tools	QUALITATIVELY NEW RISK Autonomous AI
LLMs compress expertise: protocol guidance, troubleshooting, and step-by-step planning become accessible to actors without formal training. KEY BOTTLENECK Physical lab execution Tactic knowledge, manual dexterity, and real-time troubleshooting remain hard barriers even with AI guidance. AI CAPABILITIES TESTED •Domain knowledge & protocol recall •Step-by-step planning from prompts •Physical execution (limited by format) EVALUATION APPROACH Human-AI uplift RCTs Controlled experiments measuring whether novices with LLM access complete tasks they otherwise could not. <i>Benchmarks may overestimate real-world novice uplift — strong QA performance does not translate directly to lab capability.</i>	BDTs and LLM-powered agents enable design of novel biological agents, optimised properties, and circumvention of screening measures. KEY BOTTLENECK Predictability & targeting Making engineered agents reliable, targetable, and resistant to countermeasures requires iterative experimental validation. AI CAPABILITIES TESTED •Computational design (sequences, proteins) •Bioinformatics tool use & coding •Synthesis screening evasion EVALUATION APPROACH Agentic capability benchmarks AI performance on isolated computational biology tasks requiring tool use, coding, and domain integration. <i>Strong on isolated computational sub-tasks; integration across multi-step workflows and long-horizon planning remain key bottlenecks.</i>	Misaligned or directed AI systems could independently pursue biological weapon development and deployment — removing the human bottleneck entirely. KEY BOTTLENECK Long-horizon planning & verification Reliable multi-step execution, error recovery, and operating lab equipment without human oversight. AI CAPABILITIES TESTED •End-to-end workflow completion •Autonomous tool use & error recovery •Multi-step planning across dependencies EVALUATION APPROACH THIS WORK: Modular workflow decomposition Node-by-node capability evaluation across dependency-structured threat pathways, supporting both granular and composed assessments. <i>The least measured threat model — and the one where AI capability is advancing fastest. Requires fundamentally different evaluation methodology.</i>

Figure 1. Three biosecurity threat models and their evaluation requirements.

Theory of Change

Decision-makers need evaluations that track AI capabilities in real time and forecast emerging risks, so they can act before dangerous thresholds are crossed.



BENCHMARK	FORMAT	HEADLINE RESULT
WMDP-Bio Hazardous biosecurity knowledge proxy	QA	Frontier models match or exceed expert-level performance; some biology subsets plateau below 100%, suggesting data issues.
LAB-Bench Cloning design, protocol reasoning, literature & sequence understanding	QA	Top models reach expert baseline on cloning scenarios. Some components improved >2x across model generations.
VCT Virology troubleshooting & interpretation	QA	Performance rose >4x between 2023 and 2025 ; top model outperforms expert virologists in their areas of specialty.
BioLP-Bench Lab protocol understanding & error detection	QA	Top models score above expert baseline, but best scores still <50/100 — substantial room remains.
ABC-Bench Robot programming, fragment design, synthesis screening evasion	Agentic	Top model scored 53% vs. 24% expert average ; outperformed 60-100% of experts across sub-tasks.
ABLE Protein design theory, computational tools, BAI/MS, programming	Agentic	No model completed the full workflow. Strongest model achieved perfect scores on several individual sub-tasks.
RAND Oligonucleotide design for assembly, operating synthesis equipment	Agentic	Top models achieve 9/10 full-criteria successes on eGFP; one top-performing attempt was physically validated .
Hong et al. 2026 Novice wet-lab uplift RCT on viral reverse genetics workflow (n=153)	RCT	No significant difference in end-to-end workflow completion (5.2% LLM vs 6.6% control; p=0.759). Modest ~1.4x task-level uplift.

Figure 2. Evaluation landscape: task formats, capability coverage, and expert-relative performance signals across biosecurity-relevant benchmarks.

2. Evaluation Design




Figure 3. Dependency graph for autonomous influenza virus acquisition via reverse genetics.

- Autonomous influenza virus acquisition modelled as a directed dependency graph, where nodes are discrete verifiable tasks and edges encode prerequisites.
- Grounded in public rescue protocols and interviews with PhD-level virologists.
- Policy-motivated: per-node scoring shows where capabilities are emerging, enabling evidence-based thresholds.**

PCA Oligo Design Node

- First node implemented: PCA oligonucleotide design, decomposing a target dsDNA sequence into overlapping oligos.
- Agent receives a GenBank file; operates in UK AISI Inspect (ReAct scaffold, Python/Biopython/primer3, bash, optional web search).
- Must output valid JSON of oligo IDs and sequences; strand alternation prerequisite gates score to zero if failed.




Figure 4. Evaluation setup for the PCA oligo design node.

3. Results




Figure 5. Mean PCA design scores by model. Dashed line: DNAWorks reference tool baseline.

Preliminary Findings

- Five frontier models evaluated on PCA oligo design for EGFP.
- All models except DeepSeek Reasoner reliably pass prerequisites; mapping and assembly near-ceiling.
- Primary bottleneck: thermodynamic optimization (0–32% Tm success). Models understand PCA mechanics but struggle with continuous thermodynamic constraints.
- All models below DNAWorks baseline (0.820); Gemini 3.1 Pro closest (0.729), gap driven by Tm stage.
- Preliminary: small samples (n=5), single target (EGFP), scoring rubric under development.




Figure 6. Stage-level success rates by model. Mapping and secondary structure checks are near-ceiling for most models; melting temperature (Tm) uniformity is the primary bottleneck across all models tested.

4. Practical Insights

Measurement limitations

Scores might depend on the agent scaffold; we use Inspect's built-in ReAct, but different scaffolding could change results. The scoring rubric is a development proxy that will be refined; without wet-lab validation, the mapping from scorer stages to actual PCA success rates remains uncertain. Prompting is an open question: providing too much domain context risks false positives (models appear capable because the prompt contains the reasoning), while providing too little may undercount capabilities. Sample sizes are small and only one target sequence is tested.

Policy Relevance

Modular evaluations provide early warning, showing which steps models can perform. Current DNA synthesis screening frameworks (e.g., UK guidance) focus on oligos >50–200 nt; if AI automates PCA design, it lowers the barrier to assembling genes from unscreened short oligos, potentially outpacing existing safeguards.

Safety

This project processes information that could be considered hazardous. We avoid operationally actionable protocol details: the workflow graph is presented at a high level, and evaluations use safe proxy tasks with verifiable intermediate artifacts.

This work forms part of a broader project at Measure AI and was conducted through the ERA AIXBio Fellowship,, under the research management of Christopher Stamper and with mentorship from Alex Kleinman.
Contact: zuzanna.g.matuszewska@gmail.com